


Search in Complex Environments

Jihoon Yang

Machine Learning Research Laboratory
Department of Computer Science & Engineering
Sogang University

Local search and optimization

- Previously: systematic exploration of search space
 - Solution to problem is path to goal
- Local search: evaluate and modify one or more current states; All that matters is the solution state, not the path cost to reach it
- E.g. 8-queens using complete-state formulation
 - States: 8 queens on the board, one per column
 - Actions: Move a queen to a square in the same column



Local search and optimization

- Local search is suitable for **optimization problems**
 - State space = set of “complete” configurations
 - Find best state according to some *objective function* $h(s)$
 - E.g. $h(s)$ = number of conflicts
- Many optimization problems do not fit the standard search model
- Local search = keep a single current state and move to neighboring states to improve it
- Advantages:
 - Use very little memory
 - Often find reasonable solutions in large or infinite state spaces unsuitable for systematic algorithms (e.g. solve million-queens quickly)

Hill-climbing search

- Keep a single current node and move to neighboring states to improve it
- A loop that continuously moves in the direction of increasing value
 - Chooses randomly to break ties
 - It terminates when a peak is reached where no neighbor has a higher value
 - “Like climbing Everest in thick fog with amnesia”
- Hill-climbing a.k.a. *greedy local search*, *steepest ascent/descent*

Hill-climbing search

function HILL-CLIMBING(*problem*) **return** a state that is a local maximum

current \leftarrow MAKE-NODE(problem.INITIAL-STATE)

loop do

neighbor \leftarrow a highest valued successor of *current*


if *neighbor*.VALUE \leq *current*.VALUE

then return *current*.STATE

current \leftarrow *neighbor*


Hill-climbing example

- E.g. 8-queens (complete-state formulation)
 - States: 8 queens on the board, one per column
 - Actions: Move a single queen to another square in the same column
- Heuristic function $h(n)$: the number of pairs of queens that are attacking each other (directly or indirectly)



Hill-climbing example


18	12	14	13	13	12	14	14
14	16	13	15	12	14	12	16
14	12	18	13	15	12	14	14
15	14	14	♔	13	16	13	16
♔	14	17	15	♔	14	16	16
17	♔	16	18	15	♔	15	♔
18	14	♔	15	15	14	♔	16
14	14	13	17	12	14	12	18



- The first shows a state of $h = 17$ and the h -value for each possible successor
- The second shows a local minimum in the 8-queens state space ($h = 1$)

Drawbacks

State space landscape



- Depending on initial state, can get stuck in local maxima, plateaux
- A *sideways move* can be allowed in the hope that the plateau is really a shoulder

Hill-climbing variations

- Stochastic HC
 - Random selection among the uphill moves
 - The selection probability can vary with the steepness of the uphill move
- First-choice HC
 - Generating successors randomly until a better (than the current) one is found
 - Useful when a state has many successors
- Random-restart HC
 - Restart search from random initial state
 - A reasonably good local maximum can often be found after a small number of restarts

Simulated annealing

- Escape local maxima by allowing “bad” moves
 - Idea: but gradually decrease their frequency
- T , a “temperature” controlling the probability of downward steps
- One can prove: If T decreases slowly enough, then simulated annealing search will find a global optimum with probability approaching 1
- Commonly used: $T \leftarrow cT$ with c constant close to, but smaller than, 1
- Applied to VLSI layout, airline scheduling, etc.

Simulated annealing

function SIMULATED-ANNEALING(*problem*, *schedule*) **return** a solution state
input: *problem*, a problem
 schedule, a mapping from time to temperature
local variables: T , a “temperature” controlling the probability of downward steps

$current \leftarrow \text{MAKE-NODE}(\text{problem.INITIAL-STATE})$

for $t \leftarrow 1$ **to** ∞ **do**

$T \leftarrow \text{schedule}[t]$

if $T = 0$ **then return** $current$

$next \leftarrow$ a randomly selected successor of $current$

$\Delta E \leftarrow next.VALUE - current.VALUE$

if $\Delta E > 0$ **then** $current \leftarrow next$

else $current \leftarrow next$ only with probability $e^{\Delta E / T}$

T 는 항상 양수

Local beam search

- Keep track of k states instead of one
 - Initially: k random states
 - Next: determine all successors of k states
 - If any of successors is goal \rightarrow finished
 - Else select k best from successors and repeat
- Major difference with random-restart search
 - Information is shared among k search threads
- Can suffer from lack of diversity
 - Stochastic variant: choose k successors randomly with probability proportional to state value difference

Genetic algorithms

- Inspired by the process of biological evolution
- Start with k randomly generated states/individuals (**population**)
- States are scored by evaluation function (**fitness function**)
- At each step, the most fit states are selected (*survival of the fittest*) probabilistically: used as seeds for producing the next generation population – the children or *offspring*, by means of operations such as **crossover** and **mutation**
- The process is repeated until sufficiently fit states are discovered – the best state has a score exceeding a criterion
- They have been applied successfully to a variety of learning tasks and optimization problems


Genetic algorithms

```
function GENETIC_ALGORITHM(population, FITNESS-FN) return an individual
  input: population, a set of individuals
          FITNESS-FN, a function which determines the fitness of an individual
  repeat
    new_population  $\leftarrow$  empty set
    loop for i from 1 to SIZE(population) do
      x  $\leftarrow$  RANDOM_SELECTION(population, FITNESS_FN)
      y  $\leftarrow$  RANDOM_SELECTION(population, FITNESS_FN)
      child  $\leftarrow$  REPRODUCE(x,y)
      if (small random probability) then child  $\leftarrow$  MUTATE(child)
      add child to new_population
    population  $\leftarrow$  new_population
  until some individual is fit enough or enough time has elapsed
  return the best individual in population
```

Genetic algorithms

- In basic GAs, the fundamental representation of each state is a *string* over a finite alphabet (often a string of 0s and 1s), called a **chromosome**
- The mapping depends on the problem domain, and the designers
- Genetic Operators
 - **Replication**: A chromosome is merely reproduced
 - **Crossover**: Involves the mating of two parent chromosomes to yield two new offspring by copying selected bits from each parent
 - **Mutation**: Creates a single descendant from a single parent by changing the value of a randomly chosen bit (from a 1 to 0 or vice versa)
 - Other operators: specialized to the particular representation

Genetic algorithms




Fitness

Selection


Pairs

Cross-Over


Mutation



+



=



Online search

Suppose we want to site three airports in Romania:

- 6-D state space defined by $(x_1, y_1), (x_2, y_2), (x_3, y_3)$
- objective function $f(x_1, y_1, x_2, y_2, x_3, y_3) =$
sum of squared distances from each city to nearest airport

Discretization methods turn continuous space into discrete space, e.g., **empirical gradient** considers $\pm\delta$ change in each coordinate

Gradient methods compute

$$\nabla f = \left(\frac{\partial f}{\partial x_1}, \frac{\partial f}{\partial y_1}, \frac{\partial f}{\partial x_2}, \frac{\partial f}{\partial y_2}, \frac{\partial f}{\partial x_3}, \frac{\partial f}{\partial y_3} \right)$$

to increase/reduce f , e.g., by $\mathbf{x} \leftarrow \mathbf{x} + \alpha \nabla f(\mathbf{x})$

Sometimes can solve for $\nabla f(\mathbf{x}) = 0$ exactly (e.g., with one city).

Newton–Raphson (1664, 1690) iterates $\mathbf{x} \leftarrow \mathbf{x} - \mathbf{H}_f^{-1}(\mathbf{x}) \nabla f(\mathbf{x})$

to solve $\nabla f(\mathbf{x}) = 0$, where $\mathbf{H}_{ij} = \partial^2 f / \partial x_i \partial x_j$

Online search


Offline search algorithms compute a complete solution before exec.
vs.

online search ones interleave computation and action (processing input data as they are received)

- necessary for unknown environment
(dynamic or semidynamic, and nondeterministic domains)
- ⇐ exploration problem

Online search

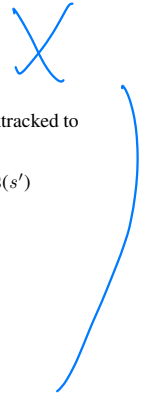
An online search agent solves problem by executing actions, rather than by pure computation (offline)



The **competitive ratio** – the total cost of the path that the agent actually travels (online cost) / that the agent would follow if it knew the search space in advance (offline cost) \Leftarrow as small as possible

Online search expands nodes in a local order, say, **DEPTHFIRST** and **HILLCLIMBING** have exactly this property

Online search



```
function ONLINE-DFS-AGENT(problem, s') returns an action
    s, a, the previous state and action, initially null
    persistent: result, a table mapping (s, a) to s', initially empty
                untried, a table mapping s to a list of untried actions
                unbacktracked, a table mapping s to a list of states never backtracked to

    if problem.IS-GOAL(s') then return stop
    if s' is a new state (not in untried) then untried[s']  $\leftarrow$  problem.ACTIONS(s')
    if s is not null then
        result[s, a]  $\leftarrow$  s'
        add s to the front of unbacktracked[s']
    if untried[s'] is empty then
        if unbacktracked[s'] is empty then return stop
        else a  $\leftarrow$  an action b such that result[s', b] = POP(unbacktracked[s'])
    else a  $\leftarrow$  POP(untried[s'])
    s  $\leftarrow$  s'
    return a
```

Figure 4.20 An online search agent that uses depth-first exploration. The agent can safely explore only in state spaces in which every action can be “undone” by some other action.